Using GPT to Predict Company Layoffs

Michael Suehle
University of Maryland,
College Park

Alex Straub
University of Maryland,
College Park

Addison Waller
University of Maryland,
College Park

Abstract

The technology job market has changed drastically over the last few years with major companies doing multiple layoffs without warning. As this trend continues we wanted to know if there is a way to predict these layoffs using alternate data, for example stock market data. To do this we have created a modeling framework utilizing transfer and ensemble learning that takes in past stock data about a company and is able to predict if a layoff has occurred or not. Further, we attempt to use time-series prediction to create future stock data to feed our model for future layoff predictions. We provide a framework that can be extended and modified to adapt to more diverse data and more advanced time series prediction. Utilizing this model, people will be able gain insight into how stock price relates to layoffs, and it opens the gate for more sophisticated models to be created to extract more accuracy for percentages and times of layoffs.

1 Introduction

Over the past couple years the technology sector has experienced a large number of mass layoffs. While the exact reason for this phenomenon is unclear, we wanted to see if there is a way to predict when the next big layoff for a company is coming. To do this we created a modeling framework that is able to predict if a company will have a layoff in a set time period. Due to time constraints, we revised our objective to be a classification problem: can we predict whether or not a layoff occurs during a three month period based on the stock data? Additionally, can forecasted stock prices help predict whether or not a layoff occurs? We decided to pursue these questions with domain adaptation and ensemble learning. First, By utilizing historical stock and layoff data, we assess if we can use time-series prediction and multi layer neural networks to predict whether or not a company will layoff people in a 90 day time frame. We test our framework with stock data from known periods with and without layoffs, and with artificially forecasted periods from our time-series forcasting.

2 Related Work

2.1 Layoff Prediction

We were only able to identify one paper that attempts to predict layoffs using machine learning. Prakash and Sakthivel use lasso regression to forecast layoffs (Prakash and Sakthivel, 2024). They do not provide the model nor results of the model's performance, but they do provide some graphs produced from exploratory analysis of the layoff data. Because of this, we could not use this as a baseline to compare against.

2.2 Incorporating Additional Data for Stock Market Forecasting

There has been some work on incorporating different types of data to improve stock market forecasting. Liao et al. incorporate company sentiment data from the news (Liao et al., 2024) and Ayyappa and Kumar incorporate political data in addition to news data to improve stock market forecasting (Ayyappa and Kumar, 2024). These studies inspired us to seek sentiment data to use for stock market prediction, however we were unable to find a suitable datasource in time for this report.

2.3 Domain Adaptation for Time Series Data

As a result of our inability to find suitable data, we looked to use domain adaptation to improve stock market predictions. We identified three different potential domain adaptation methods for time series data. Jin et al. attempt to do domain adaptation through attention sharing (Jin et al., 2022). (Cai et al., 2020) attempt to do domain adaptation with causal structure alignment between target and source data that occurs at the same time. He et al's method, RAINCOAT, attempts to do domain adaptation by training an encoder to extract frequency and time features from the source and target data and finally training a classifier on those features (He et al., 2023). We decided to use RAINCOAT because it can be applied to our problem, it achieves high performance, and its code is publicly available.

2.4 Data Transform Bagging

We also experiment with data transform bagging. Data transform bagging is an alternate version of traditional bagging where each model in the ensemble is trained on a different transformation of the same data whereas in traditional bagging, each model is trained on different samples of the same data (Brownlee, 2021).

3 Methods

Our original objective was to create a model that could predict the number of layoffs a company would have in a future time-period. To do this we attempted to do Domain Adaptation to Transform Stock Data to the domain of the layoff data. Then we wanted to input this transformed data into a forecasting model to predict the layoffs. As seen in Figure 1, the transformer mapped all the stock prices to 0. We believe this was unsuccessful due to the sparsity of our data.


Figure 1: Stock data (target data) the transformer was supposed to be able to predict compared to the actual data the transformer predicted (source data)

Due to time constraints, we revised our objective to be a classification problem: can we predict whether or not a layoff occurs during a three month period based on the stock data? Additionally, can forecasted stock prices help predict whether or not a layoff occurs? To achieve this new objective we created four different sets of data for each company that spanned three months: real stock data that contained no layoffs, real stock data that contained layoffs, chronos generated stock data with no layoffs, chronos generated stock data with layoffs. Chronos is a zero-shot time series data forecaster (Ansari et al., 2024). We used the RAINCOAT algorithm to train an encoder to extract the features of both the chronos generated data (source data) and real stock data (target data) and a classifier to predict whether a layoff occurs or not based on the extracted features. This is visualized in 2. For a baseline classifier we used the basic binary classifier and trained it on the real stock data. The classifiers would produce a 1 or 0 depending on whether there was a layoff (1) or not (0).

The next classifier we built was an ensemble of three different classifiers using a bagging method. Each of these classifiers was trained on the real stock data but the data was normalized in different ways. The first classifier was trained on z-score normalization, the second was trained on proportion to average price across periods, and the third was trained on min-max scaling. These classifiers then output a probability which represents the percentage at which they predict a layoff will happen in the set timeframe. We use this probability to create a binary classification, where when probability is greater than 0.5, it indicates there is a layoff, and less than, no layoff. Each of these classifiers performed well,


Figure 2: The architecture of RAINCOAT is as follows. A frequency feature extractor and a time features extractor are trained to produce encodings of both the target and source data. Those encodings are decoded and the reconstruction loss is back propagated through the extractors. Simultaneously, a classifier is trained on the encoding of the source data.

but to further improve performance, we implemented two different ensemble strategies. The first was to produce a confidence score by taking all the probabilities produced by the classifiers, adding them together, and dividing them by three to get an average. If that average is over 0.5, then the model predicted there was going to be a layoff, if the average is under 0.5 then the model predicted that there was not going to be a layoff. The second way to further improve performance was to take the binary conversion of the probability produced by the classifiers and perform majority voting. For example, if two classifiers output a 1 (yes layoff) and the third output a 0 (no layoff), the overall classifier would predict that there was going to be a layoff. The structure can be seen in Figure 3. Both of these bagging methods perform well, results of which will be talked about in later sections of the paper.


Figure 3: Flowchart of bagging ensemble classification method

4 Experiment

4.1 Datasets

For this project we got our layoff data from the dataset: Tech layoffs 2020 - 2024 from Kaggle (Herold, 2024). This dataset has the number of people laid off, percentage of company laid off, company size before and after layoffs, industry, and date of layoff for 1287 unique companies. It is 1672 rows total. The information in this dataset we prioritized is: Company name, Number of people laid off, Industry, and Date of layoffs. After these columns had been extracted the company names

were used to compare them to companies that exist in the Nasdaq. If they are listed then they are public companies that have stock information. Because we wanted to base our layoff predictions off stock data it is important that each company we use is public. The next step was to collect the stock information for the set of companies. Using the yfinance python library we collected each company's stock information from 03/10/2020 to 04/18/2024 which consisted of open, high, low, close, and volume. All of these data points will allow us to begin to predict the possibility of a company layoff.

We used chronos to generate our forecasted stock market data. For every layoff date in the Tech layoffs 2020-2024 dataset, we took a window of stocks one year before 45 days before the date and generated 90 days of forecasted stocks after the year of data. For our stock data to be labeled "no layoffs", we programmatically found 90 day periods of time that did not have layoffs for each company in the Tech layoffs 2020-2024 dataset and used yfinance to pull real stocks and chronos to generate forecasted stock based on the real stocks for a year 45 days before the start date of the period. For layoff stock data, we test with 260 real 90-day stock price datasets, and matching time period and stock 260 Chronos generated datasets. For no-layoffs testing we have 228 datasets for real, and 228 matching chronos generated datasets.

4.2 Evaluation

4.2.1 Baseline Classifier and RAINCOAT

We trained a baseline classifier on the real stock data and evaluated it on real stock data. To evaluate how well RAINCOAT adapts the domain of the chronos generated data to the real data, we trained another classifier using RAINCOAT and evaluated it on real stock data. To get a comparison to see if RAINCOAT improves performance using chronos generated training data, we trained a baseline classifier on the chronos generated data and evaluated it on the real stock data. We also wanted to test two more practical scenarios where we train a classifier on real data and evaluate it on chronos generated data, emulating a method for predicting if a layoff is going to occur in the future by using the classifier on projected stock prices, and training a classifier on both real and chronos generated data in the scenario where there is not enough real data, so generated data is used to increase the number of samples.

4.2.2 Ensemble Classification

We evaluate our bagging ensemble classification model across layoff and no-layoff stock datasets for both real, and chronos generated stock prices. We record each accuracy score for each of the classification models, as well as accuracy for both of our bagging ensemble methods for both real and chronos data. We also obtain 'matching' ratios, displaying how often the classifier output for the Chronos generated data agrees with the classifier output for the matching real stock data.

For an ablation study on the ensemble, we also test per-formance degradation when leaving one of the three classifiers out in the ensemble for our confidence strategy.

5 Results

5.1 RAINCOAT Dokmain Adaptation

Trained on	Tested on	Accuracy
Real	Real	0.655
Chronos	Real	0.59
Real	Chronos	0.5
Both	Real	0.51
RAINCOAT (Chronos->Real)	Real	0.57

Table 1: Accuracy scores testing the classifiers on generated and real stock data

Table 1 shows the results of our experiment. The baseline model, the classifier trained on real data and tested on real data, performed somewhat poorly, achieving an accuracy of 0.65, slightly better than random guess of 0.5. The baseline model expectedly performed significantly worse on the chronos generated data. The classifier trained on chronos generated data and evaluated on real data performed worse than the baseline model as expected, achieving an accuracy of 0.59. Performance degraded substantially when the classifier was trained on both the real and the chronos generated data. This is expected since the classifier trained on the chronos data performed poorly, but it is surprising that performance degraded as much as it did. Finally, the classifier trained with the RAINCOAT domain adaptation method performed similarly to the classifier trained on chronos data. This implies that the RAINCOAT domain adaptation did not improve the performance of the classifier.

5.2 Ensemble Classification (trained on real data)

Classifier/Ensemble	Chronos	Real
Confidence	0.5267	**0.9267**
Voting	0.5152	0.91655
MinMax	0.52305	0.7599
Open In Proportion to Average (OIPA)	0.51495	0.9075
Z-Score	0.5117	0.90285

Table 2: Accuracy scores testing the classifier(s) on generated and real stock data

For both layoff and no-layoff stock price periods, we find both the individual classifiers and the ensemble bagging output perform substantially better when using real stock data compared to the Chronos generated data (Table 2). On real data, we achieve high accuracy in all of the classifiers trained, with the Open In Proportion to Average normalization and Z-score normalization models leading the individual classifiers. Both of our ensemble methods performed better than any of the individual classifiers on real data, with >90% accuracy.

The confidence ensemble method performed the best at around 93% accuracy.

Classifier/Ensemble	Matching Ratio
Confidence	0.52155
Voting	0.5196
MinMax	0.514
Open In Proportion to Average (OIPA)	0.52465
Z-Score	0.50075

Table 3: Classification matching ratios by method between Chronos generated and real data

We find the matching ratio (the ratio of which the Chronos generated data and the matched real data predict the same label when fed into the classifier) to be quite poor. Ideally we would see a high matching ratio, since that would indicate strong model performance on future-generated Chronos data. Therefore we can not confidently say that we have created a model pipeline to predict future layoffs, since we can not rely on Chronos generated data to be representative of real data.

5.3 Ensemble Ablation

Classifier/Ensemble	Chronos	Real
Confidence (Leave MinMax Out)	0.518219	0.916565
Confidence (Leave OIPA Out)	0.53097166	0.7565789
Confidence (Leave Z-Score Out)	0.520850202	0.7598853

Table 4: Accuracy scores testing from leaving one classifier out

To bolster our results, we assess the confidence prediction when we leave a classifier out of the ensemble. For testing on real data, when leaving out MinMax, we see slightly degraded performance when compared to the complete ensemble. Omitting either the Z-score normalization model or open in proportion to the average model yielded significantly degraded performance, of around 18%. In testing on Chronos generated data, we see little or no change in confidence prediction, likely due to overall poor model performance of each of the classifiers in the ensemble for generated data.

6 Conclusion

Our results show two main conclusions. The first is that stock data generated by chronos is not sufficient for predicting whether or not a layoff occurs during the time period the data is generated for. The second is that data transformation bagging greatly improves the performance of the layoff predicting classifier when trained on real data. Although this experiment was not successful at demonstrating viability of forecasted data, it does shed light on a less common version of bagging, transform bagging. Since we achieved high accuracy on our ensemble classifier, it still may be useful for predicting a notion of layoff risk when used with recent stock data, as a positive output could indicate a layoff in the near future.

7 Future Work

Work progressing from this study can be directed in a number of ways. Utilizing more, and different types of data for ensemble classifiers to be built upon, such as industry sentiment, company sentiment, revenue, and earnings, as well as stock indicators may yield more insight into what factors contribute to layoffs. Further, utilizing a more advanced, or specified, time-series forecasting model for predicting stock price may yield significantly better performance in our classification model, enabling the use of the model to predict future layoffs. Another area of work is modifying our ensemble classification model to predict the percentage of people being laid off (vs. company size), enabling the notion of 'severity' in layoff prediction. Future work could also include modifying our model framework to predict for different time periods, including by week, or by day predictions. As layoffs continue to happen in the tech-sector, more data will become available for models like ours to use, allowing for overall better analysis and forecasting of future layoff scenarios.

References

Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, Jasper Zschiegner, Danielle C. Maddix, Hao Wang, Michael W. Mahoney, Kari Torkkola, Andrew Gordon Wilson, Michael Bohlke-Schneider, and Yuyang Wang. 2024. Chronos: Learning the language of time series.

Yalanati Ayyappa and A.P. Siva Kumar. 2024. Stock market prediction with political data analysis (sp-pda) model for handling big data. *Springer*.

Jason Brownlee. 2021. Develop a bagging ensemble with different data transformations.

Ruichu Cai, Jiawei Chen, Zijian Li, Wei Chen, Keli Zhang, Junjian Ye, Zhuozhang Li, Xiaoyan Yang, and Zhenjie Zhang. 2020. Time series domain adaptation via sparse associative structure alignment.

Huan He, Owen Queen, Teddy Koker, Consuelo Cuevas, Theodoros Tsiligkaridis, and Marinka Zitnik. 2023. Domain adaptation for time series under feature and label shifts.

Ulrike Herold. 2024. Tech layoffs 2020-2024.

Xiaoyong Jin, Youngsuk Park, Danielle C. Maddix, Hao Wang, and Yuyang Wang. 2022. Domain adaptation for time series forecasting via attention sharing.

Honglin Liao, Jiacheng Huang, and Yong Tang. 2024. Leet: stock market forecast with long-term emotional change enhanced temporal model. *PeerJ Computer Science*.

P. Prakash and V. Sakthivel. 2024. Layoffs analysis and prediction using machine learning algorithms. *Springer*.